

Preprocessing:

Histogram equalization is performed on each of the training, validation and testing images. It is a method to improve the contrast of the image using its histogram. By effectively spreading out the most frequent intensity values, areas of lower local contrast can gain a higher contrast. Since the images are colored, their color space should be first converted from RGB into one of the color spaces that separates intensity values from color components, such as YCbCr. After performing histogram equalization on the intensity plane Y, they are converted back to RGB.

Model Selection:

Multilayer convolutional neural network is chosen to fit the images. It is basically a neural network with successive convolution layer followed by a maximum pooling layer, with several fully connected layers at the end. The final layer of the network output logits of 5 corresponding to the 5 classes of flowers to be classified. A soft max function is applied to the logits to make it a probability measure sum to 1. The loss is calculated as the cross-entropy between the logits after soft max and the true one-hot label of the images.

The model is trained by stochastic gradient descent, in which at each epoch a mini batch of 100 training images are fed forward to calculate the gradients which are back propagated to update the weights of each layer. Adam optimizer is used instead of the classical stochastic gradient descent procedure. It is a more effective algorithm as it calculates an exponential moving average of the gradient and the squared gradient, and the parameters control the decay rates of these moving averages.

To select the hyperparameters of the network, such as the learning rate, depth and width of the network, sizes and strides of kernels, drop out ratio and etc., the accuracy of a validation set is monitored for every 20 iterations. If this accuracy attains a maximum, the model is saved. Early termination is implemented by introducing a patience variable. After every iteration, patience will increase by 1. If accuracy attains the maximum, patience is reset to 0. If patience exceeds 100, the training will stop.

The architecture of the network is as followings:

[224x224x3] INPUT

[56x56x96] CONV 1: 96 11x11 filters at stride 4, same padding

[28x28x96] MAX POOL 1: 2x2 filters at stride 2

[28x28x256] CONV 2: 64 5x5 filters at stride 1, same padding

[14x14x256] MAX POOL 2: 2x2 filters at stride 2

[14x14x256] CONV 3: 64 5x5 filters at stride 1, same padding

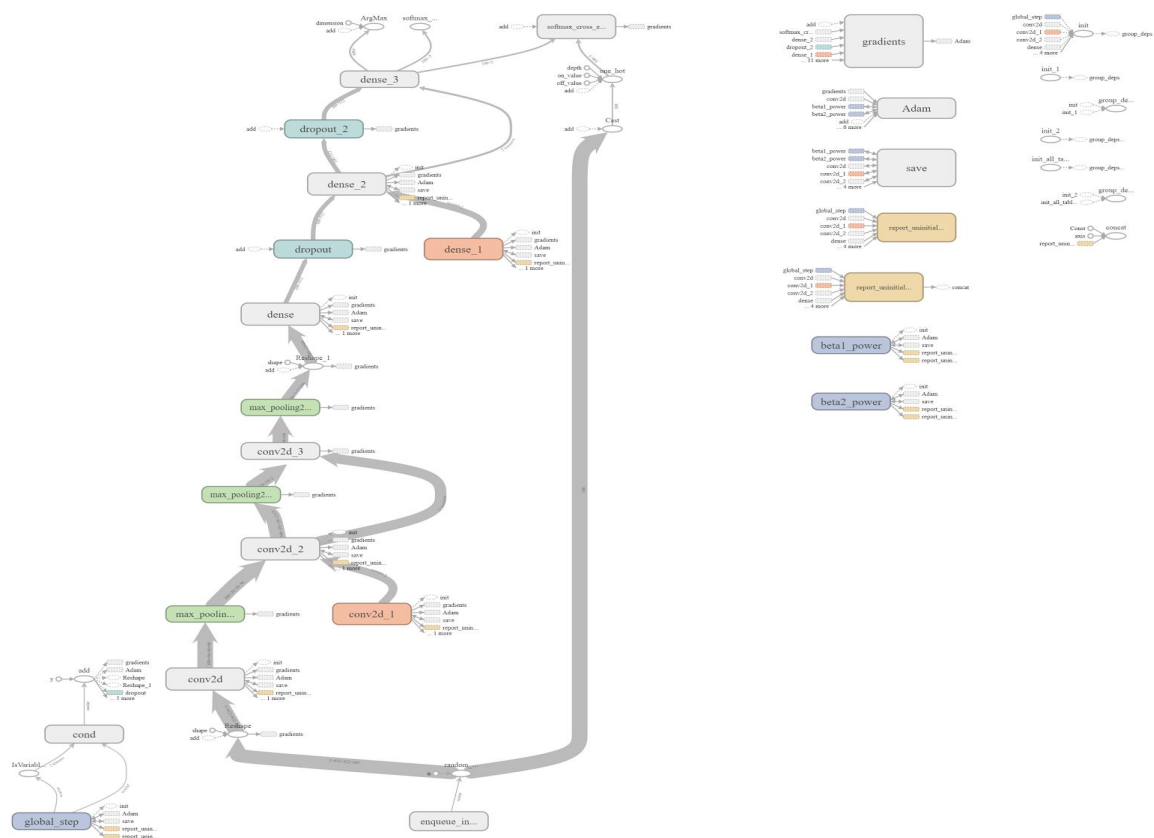
[7x7x256] MAX POOL 3: 2x2 filters at stride 2

[512] FC 4: 512 neurons


[512] FC 5: 512 neurons

[5] FC 6: 5 neurons (class scores)

Validation accuracy: 0.676364



accuracy



loss

